

Tunnels to Towers: Secure Virtual Private Networking for CCN

Abstract—Content Centric Networking (CCN) is a future Internet architecture which is envisioned as an alternative to the current IP-based model. CCN emphasizes content distribution by making contents directly addressable in a request-based information centric network. An advantage of CCN is that it has some innate privacy friendly features, such as lack of source and destination addresses in packets. However, to be considered a viable future Internet architecture, CCN must offer services for private and anonymous communication that are at least equivalent for to those present in the IP architecture. Among such, VPNs are a very popular application that enables users to send and receive data across shared or public networks as if their computing devices were directly connected to the same private network. In this work we design, implement and evaluate CCVPN, a content centric virtual private network capable of providing the same functionality of VPNs within the CCN Internet architecture. To the best of our knowledge CCVPN is the first proposal of a VPN alike service for CCNs. In addition to the CCVPN design, we also provide a security analysis and experimental performance evaluation for this new system.

Index Terms—

I. INTRODUCTION

Content-centric networking (CCN) is a type of request-based information-centric networking (ICN) architecture. In CCN, all data is named. Consumers obtain data by issuing an explicit request for the content by its name. The network is responsible for forwarding this request towards producers, based on the name, who then generate and return the content response. Since a name uniquely identifies a content response, routers may cache these packets to use in response to future requests for the same name. As a consequence, all content has an (implicit or explicit) authenticator that is used to verify the name-to-data binding. In order to prevent cache poisoning attacks, wherein a malicious producer supplies fake data in a content response that is propagated in the network, a router should never serve (a) content with an invalid authenticator or (b) cached content that it cannot verify. To enable (b), content objects with a digital signature are expected to carry the public verification key or certificate. If the authenticator is a MAC, then intermediate routers cannot verify it and should therefore not cache the content.

One negative side effect of name-based requests is that any on-path or eavesdropping adversary between a consumer and producer can learn the identity and contents of all data in transit. In traditional IP-based networks, there are generally two types of mechanisms to solve this problem: (1) anonymity networks such as Tor [?] or (2) VPNs. As tools focused on anonymity, the former help prevent linkability of packets to their requestors without always protecting the identities or content themselves.

In contrast, VPNs focus on packet confidentiality by creating a tunnel between two private networks or a consumer and single private network. All traffic over this tunnel is encrypted and thereby opaque to an eavesdropper. VPNs differ from anonymity networks such as Tor in that they are *network-layer* mechanisms that typically only introduces a single layer of encryption to protect traffic. Thus, while Tor can be used to enable VPN-like functionality, it is often far more inefficient since it operates above the network layer.

In ICN (or more specifically, CCN and NDN), ANDaNA was the first anonymity network of its kind. Similar to Tor, ANDaNA uses circuits formed from anonymizing routers (ARs) to marshal requests and responses between consumers and producers. The former onion-encrypt interests and content using the public key(s) of the target ARs. A variant of ANDaNA uses symmetric keys for packet encapsulation but suffers from linkability. Tsudik et al. [5] proposed an optimized version of the symmetric-key ANDaNA variant that did not permit linkability. To the best of our knowledge, there is no anonymity network variant for ICN architectures. Though tunneling is only useful for only a subset of ICN traffic, we believe it is a gap to be addressed for this emerging technology, for a variety of reasons. **First**, privacy continues to be an elusive property for CCN applications. Tunneling will help permit some degree of privacy within trusted AS domains from external passive eavesdroppers. **Second**, multi-hop circuits as used in ANDaNA are overkill when trying to retain privacy instead of anonymity. **Third**, end-to-end sessions such as those enabled by CCNxKE [1] and similar protocols only serve those engaged in the session. In contrast, since the threat model is different, tunneled traffic has the potential to serve any number of consumers within the same trusted domain. Thus, while tunneling may contrast the content-centric nature of data transmission in CCN, it fills a needed void for this architecture.

In this paper, we present CCVPN, a secure tunneling protocol and system design for CCN. Similar to ANDaNA, CCVPN encrypts interests and content objects between producers and consumers. In contrast, CCVPN only encrypts packets between a *single* source and destination to form a tunnel. In the standard configuration, both endpoints of the tunnel are gateways between two trusted domains. It is possible, however, for the source to be an individual consumer. In fact, the standard two-hop ANDaNA circuit is identical to a nested tunnel with the same source. Thus, CCVPN is more flexible than a general anonymity network.

XXX

II. PRELIMINARIES

This section presents an overview of the CCN architecture¹ and work related to confidentiality, privacy, and transport security. Those familiar with these topics can skip it without loss of continuity.

A. CCN Overview

In contrast to IP networks, which focus on end-host names and addresses, CCN [2], [3] centers on content by making it named, addressable, and routable within the network. A content name is a URI-like string composed of one or more variable-length name segments, each separated by a ‘/’ character. To obtain content, a user (consumer) issues a request, called an *interest* message, with the name of the desired content. This interest can be *satisfied* by either (1) a router cache or (2) the content producer. A *content object* message is returned to the consumer upon satisfaction of the interest. Moreover, name matching in CCN is exact, e.g., an interest for /edu/uci/ics/cs/fileA can only be satisfied by a content object named /edu/uci/ics/cs/fileA.

In addition to a payload, content objects include several fields. In this work, we are only interested in the following three: Name, Validation, and ExpiryTime. The Validation field is a composite of (1) validation algorithm information (e.g., the signature algorithm used, its parameters, and a link to the public verification key), and (2) validation payload (e.g., the signature). We use the term “signature” to refer to this field. ExpiryTime is an optional, producer-recommended duration for the content objects to be cached. Conversely, interest messages carry a mandatory name, optional payload, and other fields that restrict the content object response. The reader is encouraged to review [3] for a complete description of all packet fields and their semantics.

Packets are moved in the network by routers or forwarders. A forwarder is composed of at least the following two components:

- *Forwarding Interest Base (FIB)* – a table of name prefixes and corresponding outgoing interfaces. The FIB is used to route interests based on longest-prefix-matching (LPM) of their names.
- *Pending Interest Table (PIT)* – a table of outstanding (pending) interests and a set of corresponding incoming interfaces.

A forwarder may also maintain an optional *Content Store (CS)* used for content caching. The timeout for cached content is specified in the ExpiryTime field of the content header. From here on, we use the terms *CS* and *cache* interchangeably.

Forwarders use the FIB to move interests from consumers towards producers and the PIT to forward content object messages along the reverse path towards consumers. More specifically, upon receiving an interest, a router *R* first checks

its cache (if present) to see if it can satisfy this interest locally. If the content is not in the cache, *R* then consults the PIT to search for an outstanding version of the same interest. If there is a PIT match, the new incoming interface is added to the PIT entry. Otherwise, *R* forwards the interest to the next hop according to its FIB (if possible). For each forwarded interest, *R* stores some amount of state information in the PIT, including the name of the interest and the interface from which it arrived, so that content may be sent back to the consumer. When content is returned, *R* forwards it to all interfaces listed in the matching PIT entry and said entry is removed. If a router receives a content object without a matching PIT entry, the message is deemed unsolicited and subsequently discarded.

III. RELATED WORK

There are generally two classes of related work: (1) anonymity networks such as ANDaNA [4] and AC3N [5], and (2) encryption-based access control techniques. ANDaNA was developed as a proof-of-concept application-layer anonymizing network for NDN. It works by creating single-use, ephemeral, and anonymizing circuits between a consumer and producer. Each hop in this circuit uses onion decryption to decapsulate interests and onion encryption to encrypt the result. A minimum of two hops are needed to guarantee consumer and producer unlinkability. CCVPN seeks to address a different threat: privacy instead of anonymity. Thus, only a single hop, which creates a tunnel between a source and sink AS, is needed in CCVPN. Moreover, since these tunnels serve multiple traffic flows for *all* consumers within the same source, they are long-lived and persistent. This reduces the per-packet and per-flow cryptographic operations needed to move packets between the source and sink domains. AC3N [5] is an optimized version of ANDaNA that keeps per-flow state in each hop of an anonymizing circuit to specify ephemeral key identifiers for packets. This prevents intra-flow linkability while simultaneously enabling symmetric-key encryption and decryption for more efficient processing. However, since AC3N is still a per-consumer application, anonymizing circuits cannot be shared among multiple geolocated consumers in the same domain.

Content encryption seeks to solve the problem of data confidentiality rather than privacy or anonymity. This technique allows content to be disseminated throughout the network since it cannot be decrypted by adversaries without the appropriate decryption key(s). Many variations of this approach have been proposed based on general group-based encryption [6], broadcast encryption [7], [8] and proxy re-encryption [9]. Kurihara et al. [10] generalized these specialized approaches in a framework called CCN-AC, an encryption-based access control framework that shows how to use manifests to explicitly specify and enforce other encryption-based access control policies. Consumers use information in the manifest to (1) request appropriate decryption keys and (2) use them to decrypt content object(s). The NDN NBAC [11] scheme is similar to [10] in that it allows decryption keys to be flexibly specified by a data owner. However, it does this based on name engineering

¹Named-Data Networking [2] is an ICN architecture related to CCN. However, since CCNxKE was designed for ICNs that have features which are not supported by NDN (such as exact name matching), we do not focus on NDN in this work. However, CCNx could be retrofitted to work for NDN as well.

rules instead of configuration. Interest-based access control [12] is a different type of access control scheme wherein content was optionally encrypted. Access was protected by making the names of content derivable by only authorized consumers. NDN-ACE [13] is a recent access control framework for IoT environments which includes a key exchange protocol for distributing secret keys to sensors. All of these techniques, with few exceptions, use public-key cryptographic schemes to protect only the *payload* of content packets. They do not encapsulate complete packets for private transmission between a source and sink domain. The exception is CCNxKE[1], which specifies a key exchange protocol that can bootstrap secure sessions between a source a sink. CCNxKE can be used in CCVPN to establish pairwise shared secrets between tunnel endpoints. It is not a requirement for this step.

IV. CCVPN

Traditional Virtual Private Networks (VPNs) extend private networks across the Internet. They enable users to send and receive data across shared or public networks as if their computing devices were directly connected to the same private network [14]. The goal of CCVPN it to provide its users with the same functionality within the CCN Internet architecture. Therefore, the users can benefit from the features and security of a private network, even though the are not physically under the same private network.

In CCVPN, there are four main entities involved in the virtual private communication: *Consumer*, *Producer*, *Consumer Side Gateway* (G_c), and *Producer Side Gateway* (G_p). As in usual CCNs, the *Consumer* is the network node that issues an interest for a given content (e.g., file, web page, video) that it wishes to retrieve. The *Producer* is the network node which originally created such content. G_c and G_p are the edge gateways responsible for ensuring the private communication among distinct domains. As we discuss later on, G_c and G_p can actually be implemented as a single network device, but we firstly present them as separate entities for the purpose of clarity.

figure

As depicted in Fig. ??, the devices inside the *Consumer* domain form a physically interconnected private network. Conversely, the devices in the *Producer* domain also form a private network. Therefore, the goal is to create an overlay virtual network that unifies the *Consumer* and the *Producer* domains in way such that the original interests and contents are only visible to the devices inside these two domains. In other words, the interest I_e and the content C_e , that are forwarded outside these domains, should give no information about the original interest I_p and the actual content C_p .

In order to achieve such anonymous communication, G_c is introduced in the *Consumer* domain to encapsulate outgoing interests. Conversely, the G_p is responsible for decapsulating the incoming encapsulated interests and forwarding them in their original form. When the content is forwarded back in response to the interest, G_p encrypts the content before sending it outside the *Producer* domain. Finally, G_c decrypts

the received content to its original form and forwards it back towards the *Consumer*. Through the rest of this section we provide a detailed description on how such actions are implemented.

Upon the arrival of a new interest I_p , G_c checks its Forwarding Information Base (FIB) to check if the interest prefix is in the list of prefixes for VPN communication. We assume that the FIB must be pre-configured with the list of prefixes which will trigger the VPN communication. Associated with such prefixes are G_p 's name and public key (pk_e). Therefore, if the prefix of I_p is in the VPN prefixes list, G_p runs Algorithm IV to generate a new interest I_e which encapsulates the original interest I_p . Firstly, Algorithm IV generates a random symmetric key (k_r) which will be used later on to perform the *Content* Encryption/Decryption. Next, it retrieves G_p 's name and public key from the FIB. It uses the public key to encrypt the symmetric key k_r and the original interest I_p . Then it creates the new interest I_e with G_p 's name as the interest name and the generated ciphertext as payload. Since I_e has the G_p 's name it will be routed towards G_p and since the payload is signed with G_p 's public key, only G_p can retrieve the original interest I_p and the symmetric key k_r .

Algorithm 1 Interest encapsulation (runs on G_c)

input : Original interest I_p ;
output: Encapsulated interest I_e ;
 $k_r = \text{symmKeyGen}()$;
 $Gp_{name} = \text{retrieveNameFromFIB}(I_p)$
 $pk_e = \text{retrievePKFromFIB}(I_p)$
 $payload = \text{Enc}_{pk_e}(I_p || k_r)$
 $I_e = \text{createNewInterest}(Gp_{name}, payload)$
 $\text{storeToPIT}(I_e, k_r)$
return I_e ;

After the message is routed towards G_p , G_p verifies if the incoming interest name prefix matches its own name and, if it does, G_p runs Algorithm IV, using its own secret key (sk_e) to decrypt I_e 's payload, which results in the original interest I_p and the symmetric key k_r . G_p then stores I_e 's name and the symmetric key k_r in its own PIT, within the entry for the pending interest I_p , and forwards I_p . I_e 's name and k_r are stored so that they can be used later on to generate the encrypted content C_e .

Algorithm 2 Interest decapsulation (runs on G_p)

input : Encapsulated interest I_e ;
input : Private key sk_e ;
output: Original interest I_p ;
 $I_{e_{name}} = \text{getName}(I_e)$
 $cipherText = \text{getPayload}(I_e)$
 $I_p || k_r = \text{Dec}_{sk_e}(cipherText)$
 $\text{storeToPIT}(I_p, k_r, I_{e_{name}})$
return I_p ;

The original interest I_p is forwarded inside the *Producer* domain until it reaches the *Producer*. The *Producer* responds with the content C_p which is forwarded back to G_p . Upon receiving C_p , G_p fetches for C_p 's name (which is equal to I_p 's name) on its PIT, retrieving k_r and I_e 's name. Then it uses k_r to encrypt-then-MAC the real content response C_p and creates C_e which must have the same name as I_e and the encryption of C_p as payload (Algorithm IV). Since only G_c and G_p share the symmetric key k_r , only G_c will be able to decrypt C_e 's payload into C_p . Therefore nobody from outside the VPN is able to access the content nor the *Producer*'s identity.

Algorithm 3 Content encryption (runs on G_p)

input : Original content C_p ;
output: Encrypted content C_e ;
 $name = getName(C_p)$
 $k_r = retrieveKeyFromPIT(name)$
 $I_{e_name} = retrieveNameFromPIT(name)$
 $payload = EncryptThenMAC_{k_r}(C_p)$
 $C_e = createNewContent(I_{e_name}, payload)$
return C_e ;

Since C_e and I_e have the same name, C_e will be forwarded all the way back to G_c . When G_c receives C_e G_c will execute Algorithm IV. It will match C_e 's name to the pending interest I_e in its PIT, retrieving k_r . k_r can then be used to verify the integrity of the received content and to decrypt it into the actual content C_p . After that, C_p can be forwarded back to the *Consumer*. If the MAC verification fails, it means that C_e has been forged and G_c ignores it.

Algorithm 4 Content decryption (runs on G_c)

input : Encrypted content C_e ;
output: Original content C_p ;
 $C_{e_name} = getName(C_e)$
 $k_r = retrieveKeyFromPIT(C_{e_name})$
 $cipherText = getPayload(C_e)$
 $C_p = Dec(k_r, cipherText)$
if $C_p == \perp$ **then**
 /* MAC verification failed */
 return;
else
 return C_p ;
end

For clarity, we have defined a *Consumer* and a *Producer* domain. However, in reality, a single gateway can implement the functions of both G_c and G_p . Therefore, consumers and producers can exist in both the domains and interests for contents can be issued from both sides. Also, it is worth to mention that, within the created CCVPNs, content caching would work just as it works in regular CCNs, i.e., routers would be able to cache contents and respond to interests that were previously requested, enabling better resource usage and

lower communication delays. Finally, we emphasize that G_c encapsulation and decryption functions can also run inside the *Consumer* host. Conversely, G_p can be implemented within the *Producer* host. This enables its usage for one-to-one communication that would be completely anonymous to any other entity in the network.

V. CCVPN SECURITY

A. System and Security Model

Our security analysis considers the worst case scenario, in which the consumer issues an interest that has not been cached in any of the routers from the consumer nor producer domains. Therefore, the interest must traverse the whole network until reaching the content producer. Conversely, the content needs to be forwarded all the way back though the same path satisfying the routers' PIT entries. A system that is secure in this worst case is also secure when the content has been previously cached in some router along the way.

Adversary Goals and Capabilities. The goal of the adversary is either 1) to retrieve any information associated with the original issued interest I_p and/or the original content C_p or 2) to retrieve the identities of the consumer and/or producer, thus violating the anonymity in the communication. An adversary is also considered successful if it is able to impersonate the content producer by faking a content response. Since the original interest and the original content are visible inside the consumer and producer domains, we here assume that the adversary cannot eavesdrop or compromise hosts inside such domains. It is also worth to emphasize that, by the conceptual design of the VPN, communication is anonymous to hosts from outside the VPN but not to hosts that belong to that VPN. Anyhow, one-to-one anonymous communication is still achievable using the VPN paradigm if the communication must be anonymous to any other host in the network². This also means that, as components of the producer and consumer domains, we assume that the gateways G_c and G_p cannot be compromised.

In our model the adversary is allowed to perform the following actions (when outside the consumer and producer domains):

- **Eavesdrop and replay traffic:** An adversary can eavesdrop on a link learning among other things the packet contents and traffic patterns.
- **Deploy compromised routers or compromise existing routers:** An adversary is capable of deploying a compromised router or compromised an existent router outside the VPN domains. By doing so the adversary becomes capable of maliciously injecting, delaying, altering, and dropping traffic. In addition, when an existing router is compromised, the adversary learns all of the private

²In the case that communication should be anonymous even to nodes inside the consumer and producer domains, G_c and G_p could be ran as processes within the consumer and producer hosts, respectively. This way, packets exchanged between them would remain anonymous to any other node in the network

information contained in that router, such as private keys and cached content.

- **Deploy compromised caches:** As a consequence of the ability to compromise and deploy compromised routers, an adversary is also able to deploy compromised caches. This includes monitoring the routers' cache contents and replying with corrupted or fake data.

B. Additional Considerations

Gateway-to-Gateway Authentication: At this point one might have noticed that, in the CCVPN design, any host that possesses the producer -side gateway (G_p) public-key is able to initiate an anonymous communication link with G_p . In other words, the design does not include any authentication between the consumer-side gateway G_c and G_p . Indeed, authentication is not required in CCVPN because is in not needed in some application scenarios.

Suppose, for instance, that a content provider offers its contents to any host in the Internet but it also wants such hosts to be able to anonymously request and receive such content. In this case, since any host in the Internet should be able to requests the contents, there is no need for the consumer-side G_c to authenticate it self to G_p .

Another use case for CCVPN is the traditional VPN use case, in which two physically separated local networks (for instances to offices of the same company in different countries) should virtually behave as a private local network. In that case, it is necessary to prevent that a given G_c , which is not part of that corporation's network to connect to G_p and become part of the VPN. To that purpose, standard host-to-host CCN authentication mechanisms such as [] can be used to authenticate G_p and G_c to each other. This must be performed before any VPN communication to make sure that only the appropriate parties are communicating under the CCVPN architecture. We discuss the evaluation and specification of gateway-to-gateway authentication protocols as future work (see Sec. VIII).

Denial of Service: Since CCVPN gateways are connected to the public network they are clearly susceptible to DoS attacks. In the CCVPN architecture, both G_c and G_p are susceptible to DoS attacks. A DoS attack on G_p would consist in flooding it with several fake encapsulated interests. Conversely, a DoS attack on G_c would basically consist of flooding it with an enormous amount of encrypted content packets. In the case of G_p , such attacks are specially harmful since the interest decapsulation involves a public-key decryption operation.

C. Security Analysis

TODO: Introductory text about the sec analysis

Definition 5.1: An interest encapsulation algorithm $Encapsulate(I_p)$ is an indistinguishable interest encapsulation iff, given any two interests I_p^1 and I_p^2 , chosen by the adversary, and a randomly selected bit b , the adversary has only $1/2 + \epsilon$ probability of guessing the value of the bit b when given $I_e^b = Encapsulate(I_p^b)$. Where ϵ is a negligible factor with regard to the security parameter k .

Claim 5.1: Let $Encapsulate_{pk}(I_p)$ denote the interest encapsulation routine described in Algorithm IV. If Enc_{pk} , used to construct $Encapsulate_{pk}(I_p)$, is a CPA-secure public-key encryption scheme then $Encapsulate_{pk}(I_p)$ is an indistinguishable interest encapsulation algorithm.

Proof– Suppose that Claim 5.1 is false. Then there exists a polynomial adversary Adv capable of guessing the bit b of Definition 5.1 with non-negligible advantage, when given $I_e^b = Encapsulate(I_p^b)$ with $b \leftarrow \{0, 1\}$ chosen at random. We show that if such adversary exists he can be used to construct a polynomial adversary Adv_{CPA} which breaks the CPA-security of Enc_{pk} . Adv_{CPA} plays the CPA-security game with a challenger sending him two messages m^0 and m^1 . Following the CPA-security game, the challenger randomly chooses a value for the bit $b' \leftarrow \{0, 1\}$ and gives back $C = Enc_{pk}(m^{b'})$ to Adv_{CPA} . To break the CPA-security Adv_{CPA} must be able to guess the value of the bit b' with non-negligible advantage. To that purpose Adv_{CPA} can query the challenger for the encryptions of m^0 and m^1 ($c^0 = Enc_{pk}(m^0)$ and $c^1 = Enc_{pk}(m^1)$) and then construct two interests $I_e^0 = createNewInterest(Gp_{name}, c^0)$ and $I_e^1 = createNewInterest(Gp_{name}, c^1)$, using the same $createNewInterest$ function used by algorithm IV, which is public (notice that Gp_{name} is also public). Finally, Adv_{CPA} gives I_e^0 and I_e^1 as input to Adv and outputs whatever Adv outputs. Since under our assumption Adv guesses the bit b with non-negligible advantage, then Adv_{CPA} breaks the CPA-security of Enc_{pk} . But this violates the hypothesis of Claim 5.1 and, therefore, such Adv cannot exist.

Definition 5.2: A content encapsulation algorithm $Encapsulate(C_p)$ is an indistinguishable content encapsulation iff, given any two contents C_p^1 and C_p^2 , chosen by the adversary, and a randomly selected bit b , the adversary has only $1/2 + \epsilon$ probability of guessing the value of the bit b when given $C_e^b = Encapsulate(C_p^b)$. Where ϵ is a negligible factor with regard to the security parameter k .

Claim 5.2: Let $ContentEnc_{kr}(C_p)$ denote the content encapsulation routine described in Algorithm IV. If $EncryptThenMAC_{kr}$, used to construct $ContentEnc_{sk}$, is a CCA-secure symmetric-key encryption scheme, then:

- 1) $ContentEnc_{kr}(C_p)$ is an indistinguishable content encapsulation algorithm;
- 2) An adversary has only negligible probability of generating a valid fake encapsulated content I'_c

Proof– TODO

VI. PERFORMANCE ANALYSIS

In this section we analyze and discuss the overhead of the CCVPN design with respect to the additional processing time and state consumption needed to handle traffic.

A. State Consumption

The CCVPN design has an immediate impact on the FIB and PIT size of a gateway. (The content store size remains unaffected since only decapsulated content objects are ever cached.) Let F_S be the total size of a standard forwarder FIB

in terms of bytes and N_F be the number of entries in the FIB. For simplicity, we will assume that each name prefix in the FIB has a constant size of 64B. In practice we expect this to be a comfortable upper bound. Thus, $F_S = N_F s$, where s is the size of each FIB entry. Here, s includes a name prefix (of size 64B) and a bit vector that identifies the matching links for the interface. We assume that a gateway has 128 links which, again, is a comfortable upper bound. Therefore, $F_S = 80N_FB$. Now consider the FIB size F_G for a CCVPN gateway. Some entries in these FIBs will point to “private” prefixes, i.e., other domains, and therefore have a larger size to account for the corresponding prefix and key material that must be stored. For both public- and symmetric-key encryption, the key size is the same: 32B [?]. Therefore, by taking into account two both the FIB entry prefix key, translation prefix, encryption key, and corresponding bit vector, the total size of one “private” FIB entry will be 176B, meaning that $F_G = 176N_FB$. By comparing F_S to F_G , we see that, in the worst case, the CCVPN FIB is at most $F_G/F_S = 176/80 = 2.2$ times larger than the standard FIB. In practice, however, we expect this to be much smaller, since the fraction of public to private FIB entries in a gateway will be non-zero.

We will now apply the same analysis to the PIT size. A standard PIT entry includes a complete name and ingress bit vector. (They may also include the optional `KeyId` and `ContentId`, but since they are included in the gateway PIT as well we omit them from this analysis.) A gateway PIT entry will contain the same elements of a standard PIT entry but also a symmetric encryption key (32B), nonce (12B), and an encapsulation name (64B + 32B). The encapsulation name is the name of an encapsulated interest and includes an additional 32B `PayloadID` segment to identify the encapsulated value in the payload. Let P_S and P_G be the sizes of the standard and gateway PIT, respectively, and let N_P be the number of PIT entries in one such table. Based on the above discussion, and assuming again that a name is at most 64B, a standard PIT entry is of the size 80B. In contrast, a gateway PIT entry is of size 204B. Therefore, in the worst case, the CCVPN PIT will be at most $P_G/P_S = 204/80 = 2.55$ larger than the standard PIT. Assuming a steady state size of approximately $1e^5$ entries [15], this means that the PIT will be 20.4MB, which is well within the capacity of modern memory systems.

B. Processing Overhead

In terms of processing overhead, the gateway adds a number of new steps to the data path of a packet. The main computational burdens are packet encapsulation and decapsulation. In the public-key variant of CCVPN, interests are processed using public-key encryption, whereas content is always processed using symmetric-key encryption. Let $T_E^P(n)$ and $T_D^P(n)$ be the time to encrypt and decrypt nB of data using a suitable public-key encryption scheme. Similarly, let $T_E^S(n)$ and $T_D^S(n)$ be the time to encrypt and decrypt nB of data using a symmetric-key encryption scheme. Then, the latency in a single interest-content exchange is increased by $T = T_E^P(n_I) + T_D^P(n_I) + T_E^S(n_C) + T_D^S(n_C)$, where n_I and

n_C are the original interest and content sizes, respectively. As a rough estimate, [16] lists the cost of AES-GCM to be $2.946\mu s$ for setup followed by 102MiB/second Intel Core 2 1.83 GHz processor under Windows Vista in 32-bit mode (with AES ISA support). For packets that are at most 1500B, the total processing time is roughly $17\mu s$. Moreover, The public-key encryption and decryption operations will always be at least as expensive, so the total latency is increased by at least $T = 4 \times 17\mu s = 68\mu s$. In comparison to the network latency for a single packet this may not be noticable, but for a steady arrival state of approximately $1e^5$, this would lead to an instable system that would quickly overflow. (This is because $65\mu s \times 1e^5 = 6.8s$.) Therefore, there is an upper bound on the number of private packets a gateway can process per second. This bound is entirely dependent on the system configuration and network conditions.

Another performance deficiency comes from the fact that gateways cannot process packets without allocating memory. Specifically, since each packet requires either an encryption or decryption, which cannot be done entirely in-place, the gateway must allocate some amount of memory for every processed packet. This overhead can outweigh the cryptographic computations if the packet arrival rate is high enough. Therefore, when implementing CCVPN, special care must be taken to ensure that all cryptographic operations are performed in-place where possible.

VII. IMPLEMENTATION AND EXPERIMENTS

A. Experimental Methodology

In this section we empirically evaluate the CCVPN design paying special attention to the metrics that were earlier discussed in Sec. VI, i.e., processing overhead, network throughput, and state consumption. In our evaluation we consider the two versions of CCVPN: public key version, and symmetric key version. We here recall that the symmetric key version relies on the assumption that a secure key agreement protocol is performed between the domain’s gateways prior to the CCVPN protocol execution.

Our testbed network consists of a butterfly topology, in which the consumers’ side and the producers’ side gateways are directly interconnected. N producers are connected to the producers’ domain gateway and M consumers are connected to the consumers’ domain gateway (see Fig. ??).

To investigate the processing overhead we measure the average time demand for computing the interests’ encapsulation (in public and symmetric key versions), interest decapsulation, content encryption, and content decryption for different content packet sizes (1024, 4096, 16384, and 65536 bytes). We also measure the state consumption for these same four functions.

To compute the overall network throughput we measure the average data-rate for transmissions of 1 to 1,000,000 different interests issues per consumer. We also vary the number of consumers and producers from 1 to 10 of each. Finally, in addition to the network throughput, we also exhibit the total transmission delay for each of the experiments.

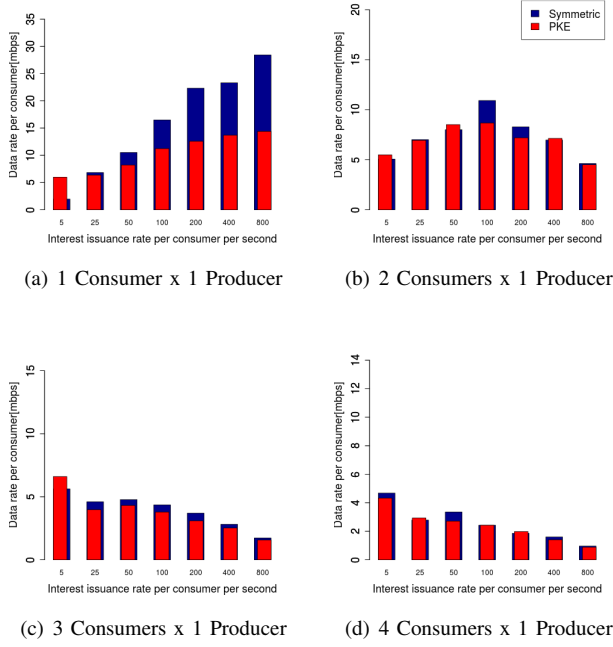


Fig. 2. Throughput per consumer. 1 Producer X N consumers. Content packet size: 1024B

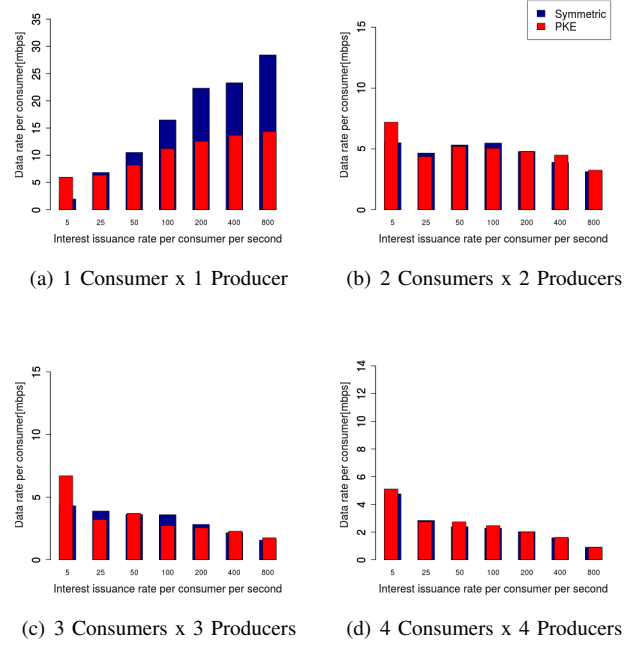


Fig. 3. Throughput per consumer. N Producers X N consumers. Content packet size: 1024B

B. Results

TODO: Add results discussion

TABLE I
INTEREST ENCAPSULATION PROCESSING TIMES

Encapsulation mode	Encapsulation	Decapsulation
Public Key	444 μ s	449 μ s
Symmetric Key	TODO	TODO

TABLE II
CONTENT ENCRYPTION AND DECRYPTION TIMES FOR DIFFERENT PAYLOAD SIZES

Packet size	Encryption	Decryption
1024B	125 μ s	193 μ s
4096B	141 μ s	220 μ s
16384B	220 μ s	367 μ s
65536B	519 μ s	702 μ s

VIII. CONCLUSION

TODO

REFERENCES

- [1] M. Mosko, E. Uzun, and C. Wood, "CCNx Key Exchange Protocol Version 1.0," Tech. Rep. draft-wood-ccnxkeyexchange-01, May 2016. [Online]. Available: <https://raw.githubusercontent.com/PARC/ccnx-keyexchange-rfc/master/draft-wood-icnrg-ccnxkeyexchange-01.txt>
- [2] V. Jacobson *et al.*, "Networking named content," in *CoNext*, 2009.
- [3] M. Mosko, I. Solis, and C. Wood, "CCNx semantics," *IRTF Draft, Palo Alto Research Center, Inc.*, 2016.
- [4] S. DiBenedetto, P. Gasti, G. Tsudik, and E. Uzun, "Andana: Anonymous named data networking application," *arXiv preprint arXiv:1112.2205*, 2011.

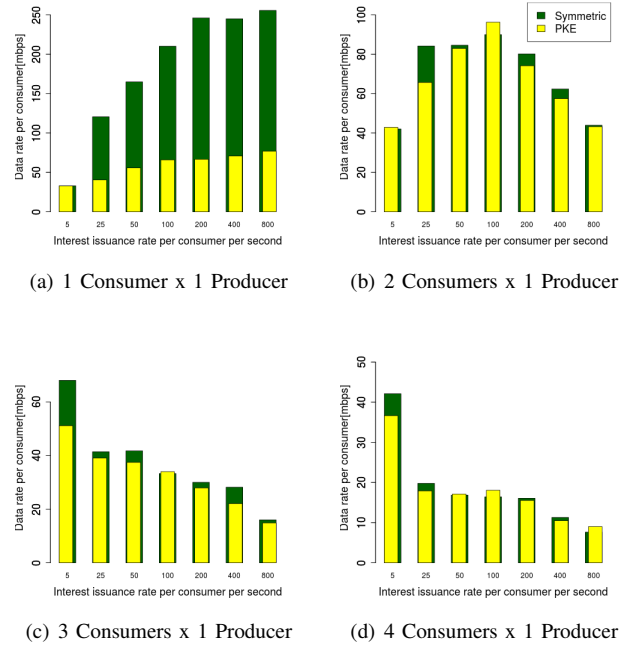


Fig. 4. Throughput per consumer. 1 Producer X N consumers. Content packet size: 10240B

- [5] G. Tsudik, E. Uzun, and C. A. Wood, "Ac3n: Anonymous communication in content-centric networking," in *2016 13th IEEE Annual Consumer Communications & Networking Conference (CCNC)*. IEEE, 2016, pp. 988–991.
- [6] D. K. Smetters, P. Golle, and J. D. Thornton, "CCNx access control specifications," PARC, Tech. Rep., Jul. 2010.

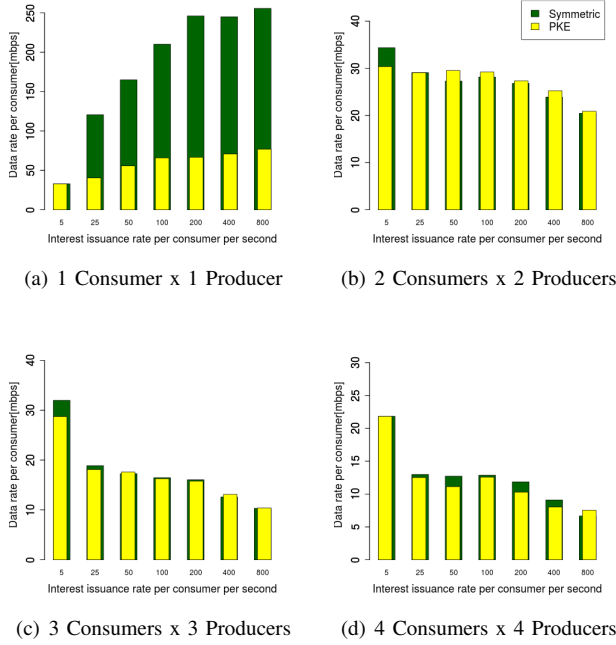


Fig. 5. Throughput per consumer. N Producers X N consumers. Content packet size: 10240B

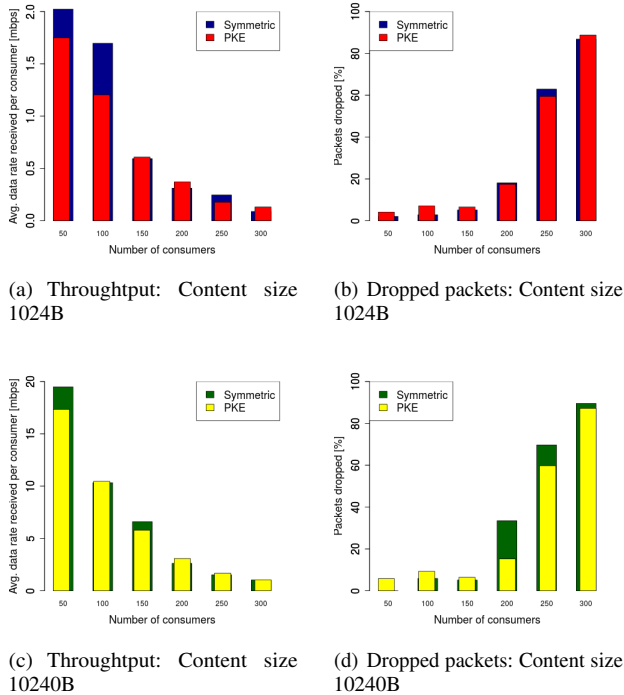


Fig. 6. Throughput per consumer with massive number of consumers/producers. Each consumer issues 5 interests

- [9] C. A. Wood and E. Uzun, "Flexible end-to-end content security in CCN," in *CCNC*, 2014.
- [10] J. Kurihara, C. Wood, and E. Uzun, "An encryption-based access control framework for content-centric networking," *IFIP*, 2015.
- [11] Y. Yu, A. Afanasyev, and L. Zhang, "Name-based access control," *Named Data Networking Project, Technical Report NDN-0034*, 2015.
- [12] C. Ghali, M. A. Schlosberg, G. Tsudik, and C. A. Wood, "Interest-based access control for content centric networks," in *International Conference on Information-Centric Networking*. ACM, 2015.
- [13] W. Shang *et al.*, "Ndn-ace: Access control for constrained environments over named data networking."
- [14] S. Khanvilkar and A. Khokhar, "Virtual private networks: an overview with performance evaluation," *IEEE Communications Magazine*, vol. 42, no. 10, pp. 146–154, 2004.
- [15] G. Carofiglio, M. Gallo, L. Muscariello, and D. Perino, "Pending interest table sizing in named data networking," in *Proceedings of the 2nd International Conference on Information-Centric Networking*. ACM, 2015, pp. 49–58.
- [16] "Crypto++ 5.6.0 Benchmarks," <https://www.cryptopp.com/benchmarks.html>, accessed: 2016-11-21.

- [7] S. Misra, R. Tourani, and N. E. Majd, "Secure content delivery in information-centric networks: Design, implementation, and analyses," in *ICN*, 2013.
- [8] M. Ion, J. Zhang, and E. M. Schooler, "Toward content-centric privacy in ICN: Attribute-based encryption and routing," in *ICN*, 2013.