CCVPN: Content Centric Virtual Private Networking

Abstract—

Index Terms—

I. Introduction

II. Related Work

- [?]
- [?]
- ...

III. CCVPN

Traditional Virtual Private Networks (VPNs) extend private networks across the Internet. They enable users to send and receive data across shared or public networks as if their computing devices were directly connected to the same private network [?]. The goal of CCVPN it to provide its users with the same functionality within the CCN Internet architecture. Therefore, the users can benefit from the features and security of a private network, even though the are not physically under the same private network.

In CCVPN, there are four main entities involved in the virtual private communication: *Consumer*, *Producer*, *Consumer Side Gateway* (G_c), and *Producer Side Gateway* (G_p). As in usual CCNs, the *Consumer* is the network node that issues an interest for a given content (e.g., file, web page, video) that it wishes to retrieve. The Producer is the network node which originally created such content. G_c and G_p are the edge gateways responsible for ensuring the private communication among distinct domains. As we discuss later on, G_c and G_p can actually be implemented as a single network device, but we firstly present them as separate entities for the purpose of clarity.

As depicted in Fig. **??**, the devices inside the *Consumer* domain form a physically interconnected private network. Conversely, the devices in the *Producer* domain also form a private network. Therefore, the goal is to create an overlay virtual network that unifies the *Consumer* and the *Producer* domains in way such that the original interests and contents are only visible to the devices inside these two domains. In other words, the interest I_e and the content C_e, that are forwarded outside these domains, should give no information about the original interest I_p and the actual content C_p.

In order to achieve such anonymous communication, G_c is introduced in the *Consumer* domain to encapsulate outgoing interests. Conversely, the G_p is responsible for decapsulating the incoming encapsulated interests and forwarding them in their original form. When the content is forwarded back in response to the interest, G_p encrypts the content before sending it outside the *Producer* domain. Finally, G_c decrypts the received content to its original form and forwards it back towards the *Consumer*. Through the rest of this section we provide a detailed description on how such actions are implemented.

Upon the arrival of a new interest I_p, G_c checks its Forwarding Information Base (FIB) to check if the interest prefix is in the list of prefixes for VPN communication. We assume that the FIB must be pre-configured with the list of prefixes which will trigger the VPN communication. Associated with such prefixes are G_p's name and public key (pk_e). Therefore, if the prefix of I_p is in the VPN prefixes list, G_P runs Algorithm **??** to generate a new interest I_e which encapsulates the original interest I_p. Firstly, Algorithm **??** generates a random symmetric key (k_r) which will be used later on to perform the *Content* Encryption/Decryption. Next, it retrieves G_p's name and public key from the FIB. It uses the public key to encrypt the symmetric key k_r and the original interest I_p. Then it creates the new interest I_e with G_p's name as the interest name and the generated ciphertext as payload. Since I_e has the G_p's name it will be routed towards G_p and since the payload is singed with G_p's public key, only G_p can retrieve the original interest I_p and the symmetric key k_r.

---

input : Original interest I_p;
output: Encapsulated interest I_e;
k_r = symmKeyGen();
Gp_{name} = retrieveNameFromFIB(I_p)
pk_e = retrievePKFromFIB(I_p)
$payload = Enc_{pk_e}(I_p || k_r)$
I_e = createNewInterest(Gp_{name}, $payload$)
storeToPIT(I_e,k_r)
return I_e;

Algorithm 1: Interest encapsulation (runs on G_c)

After the message is routed towards G_p, G_p verifies if the incoming interest name prefix matches its own name and, if it does, G_p runs Algorithm **??**, using its own secret key (sk_e) to decrypt I_e's payload, which results in the original interest I_p and the symmetric key k_r. G_p then stores I_e's name and the symmetric key k_r in its own PIT, within the entry for the pending interest I_p, and forwards I_p. I_e's name and k_r are stored so that they can be used later on to generate the encrypted content C_e.

The original interest I_p is forwarded inside the *Producer* domain until it reaches the *Producer*. The *Producer* responds with the content C_p which is forwarded back to G_p. Upon
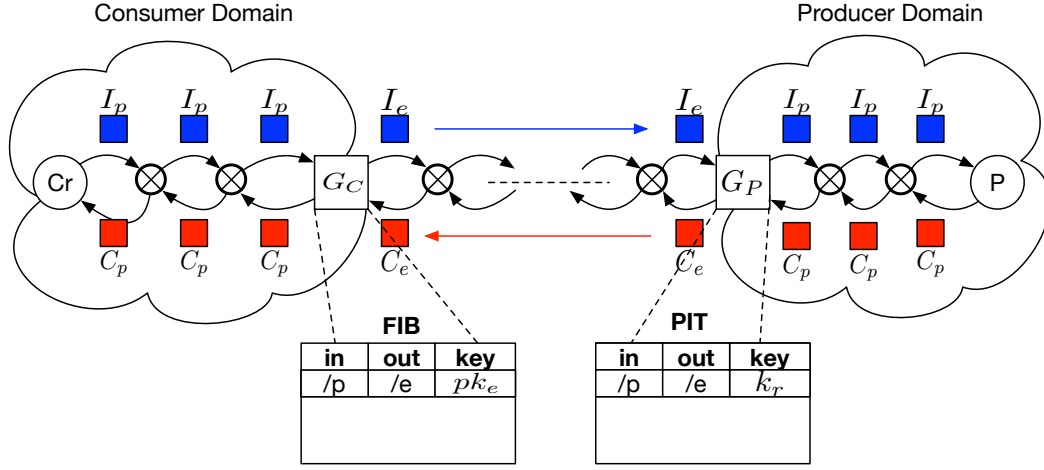
Fig. 1. CCVPN connectivity architecture

input : Encapsulated interest I_e;
input : Private key sk_e;
output: Original interest I_p;
Ie_{name} = getName(I_e)
$cipherText$ = getPayload(I_e)
$I_p || k_r$ = $Dec_{sk_e}(cipherText)$
storeToPIT(I_p,k_r,Ie_{name})
return I_p;
    Algorithm 2: Interest decapsulation (runs on G_p)

input : Original content C_p;
output: Encrypted content C_e;
$name$ = getName(C_p)
k_r = retrieveKeyFromPIT($name$)
Ie_{name} = retrieveNameFromPIT($name$)
$payload$ = EncryptThenMAC(k_r,C_p)
C_e = createNewContent(Ie_{name},$payload$)
return C_e;
    Algorithm 3: Content encryption (runs on G_p)

receiving C_p, G_p fetches for C_p's name (which is equal to I_p's name) on its PIT, retrieving k_r and Ie's name. Then it uses k_r to encrypt-then-MAC the real content response C_p and creates C_e which must have the same name as I_e and the encryption of C_p as payload (Algorithm **??**). Since only G_c and G_p share the symmetric key k_r, only G_c will be able to decrypt C_e's payload into C_p. Therefore nobody from outside the VPN is able to access the content nor the Producer's identity.

Since C_e and I_e have the same name, C_e will be forwarded all the way back to G_c. When G_c receives C_e G_c will execute Algorithm **??**. It will match C_e's name to the pending interest I_e in its PIT, retrieving k_r. k_r can then be used to verify the integrity of the received content and to decrypt it into the actual content C_p. After that, C_p can be forwarded back to the *Consumer*. If the MAC verification fails, it means that C_e has been forged and G_c ignores it.

For clarity, we have defined a *Consumer* and a *Producer*

input : Encrypted content C_e;
output: Original content C_p;
Ce_{name} = getName(C_e)
k_r = retrieveKeyFromPIT(Ce_{name})
$cipherText$ = getPayload(C_e)
C_p = Dec(k_r,$cipherText$)
if C_p == $\bot$ **then**
   | /* MAC verification failed   */
   | **return**;
else
   | **return** C_p;
end
    Algorithm 4: Content decryption (runs on G_c)

domain. However, in reality, a single gateway can implement the functions of both G_c and G_p. Therefore, consumers and producers can exist in both the domains and interests for contents can be issued from both sides. Also, it is worth to mention that, within the created CCVPNs, content caching would work just as it works in regular CCNs, i.e., routers would be able to cache contents and respond to interests that were previously requested, enabling better resource usage and lower communication delays. Finally, we emphasize that G_c encapsulation and decryption functions can also run inside the Consumer host. Conversely, G_p can be implemented within the *Producer* host. This enables its usage for one-to-one communication that would be completely anonymous to any other entity in the network.

IV. DISCUSSION AND ANALYSIS

In this section we analyze and discuss the overhead of the CCVPN design with respect to the additional processing time and state consumption needed to handle traffic.

A. State Consumption

The CCVPN design has an immediate impact on the FIB and PIT size of a gateway. (The content store size remains unaffected since only decapsulated content objects are ever cached.) Let F_S be the total size of a standard forwarder FIB in terms of bytes and N_F be the number of entries in the FIB. For simplicity, we will assume that each name prefix in the FIB has a constant size of 64B. In practice we expect this to be a comfortable upper bound. Thus, $F_S = N_F s$, where s is the size of each FIB entry. Here, s includes a name prefix (of size 64B) and a bit vector that identifies the matching links for the interface. We assume that a gateway has 128 links which, again, is a comfortable upper bound. Therefore, $F_S = 80 N_F$B. Now consider the FIB size F_G for a CCVPN gateway. Some entries in these FIBs will point to "private" prefixes, i.e., other domains, and therefore have a larger size to account for the corresponding prefix and key material that must be stored. For both public- and symmetric-key encryption, the key size is the same: 32B [**?**]. Therefore, by taking into account two both the FIB entry prefix key, translation prefix, encryption key, and corresponding bit vector, the total size of one "private" FIB entry will be 176B, meaning that $F_G = 176 N_F$B. By comparing F_S to F_G, we see that, in the worst case, the CCVPN FIB is at most $F_G/F_S = 176/80 = 2.2$ times larger than the standard FIB. In practice, however, we expect this to be much smaller, since the fraction of public to private FIB entries in a gateway will be non-zero.

We will now apply the same analysis to the PIT size. A standard PIT entry includes a complete name and ingress bit vector. (They may also include the optional `KeyId` and `ContentId`, but since they are included in the gateway PIT as well we omit them from this analysis.) A gateway PIT entry will contain the same elements of a standard PIT entry but also a symmetric encryption key (32B), nonce (12B), and an encapsulation name (64B + 32B). The encapsulation name is the name of an encapsulated interest and includes an additional

32B `PayloadID` segment to identify the encapsulated value in the payload. Let P_S and P_G be the sizes of the standard and gateway PIT, respectively, and let N_P be the number of PIT entries in one such table. Based on the above discussion, and assuming again that a name is at most 64B, a standard PIT entry is of the size 80B. In contrast, a gateway PIT entry is of size 204B. Therefore, in the worst case, the CCVPN PIT will be at most $P_G/P_S = 204/80 = 2.55$B larger than the standard PIT. Assuming a steady state size of approximately $1e^5$ entries [**?**], this means that the PIT will be 20.4MB, which is well within the capacity of modern memory systems.

B. Processing Overhead

In terms of processing overhead, the gateway adds a number of new steps to the data path of a packet. The main computational burdens are packet encapsulation and decapsulation. In the public-key variant of CCVPN, interests are processed using public-key encryption, whereas content is always processed using symmetric-key encryption. Let $T_E^P(n)$ and $T_D^P(n)$ be the time to encrypt and decrypt nB of data using a suitable public-key encryption scheme. Similarly, let $T_E^S(n)$ and $T_D^S(n)$ be the time to encrypt and decrypt nB of data using a symmetric-key encryption scheme. Then, the latency in a single interest-content exchange is increased by $T = T_E^P(n_I) + T_D^P(n_I) + T_E^S(n_C) + T_D^S(n_C)$, where n_I and n_C are the original interest and content sizes, respectively. As a rough estimate, [**?**] lists the cost of AES-GCM to be $2.946\mu s$ for setup followed by 102MiB/second Intel Core 2 1.83 GHz processor under Windows Vista in 32-bit mode (with AES ISA support). For packets that are at most 1500B, the total processing time is roughly $17\mu s$. Moreover, The public-key encryption and decryption operations will always be at least as expensive, so the total latency is increased by at least $T = 4 \times 17\mu s = 68\mu s$. In comparison to the network latency for a single packet this may not be noticable, but for a steady arrival state of approximatey $1e^5$, this would lead to an instable system that would quickly overflow. (This is because $65\mu s \times 1e^5 = 6.8s$.) Therefore, there is an upper bound on the number of private packets a gateway can process per second. This bound is entirely dependent on the system configuration and network conditions.

Another performance deficiency comes from the fact that gateways cannot process packets without allocating memory. Specifically, since each packet requires either an encryption or decryption, which cannot be done entirely in-place, the gateway must allocate some amount of memory for every processed packet. This overhead can outweigh the cryptographic computations if the packet arrival rate is high enough. Therefore, when implementing CCVPN, special care must be taken to ensure that all cryptographic operations are performed in-place where possible.

V. EXPERIMENTS

A. Experimental Methodology

B. Results

VI. CONCLUSION